

CS 240 – Data Structures and Data Management

Module 11: External Memory

Assignment Project Exam Help

Mark Petrick Olga Veksler

<https://tutorcs.com>

Based on lecture notes by many previous cs240 instructors

David R. Cheriton School of Computer Science, University of Waterloo

WeChat: [estutorcs](#)

Winter 2020

References: Goodrich & Tamassia 20.1-20.3, Sedgewick 16.4

version 2020-03-28 10:50

Assignment Project Exam Help

1 External Memory

- Motivation
- External sorting
- External Dictionaries
- 2-4 Trees
- a - b -Trees
- B-Trees
- Extendible Hashing

<https://tutorcs.com>

WeChat: cstutorcs

Assignment Project Exam Help

1 External Memory

- Motivation
- External sorting
- External Disks/Shares
- 2-4 Trees
- a - b -Trees
- B-Trees
- Extendible Hashing

<https://tutorcs.com>

WeChat: cstutorcs

Different levels of memory

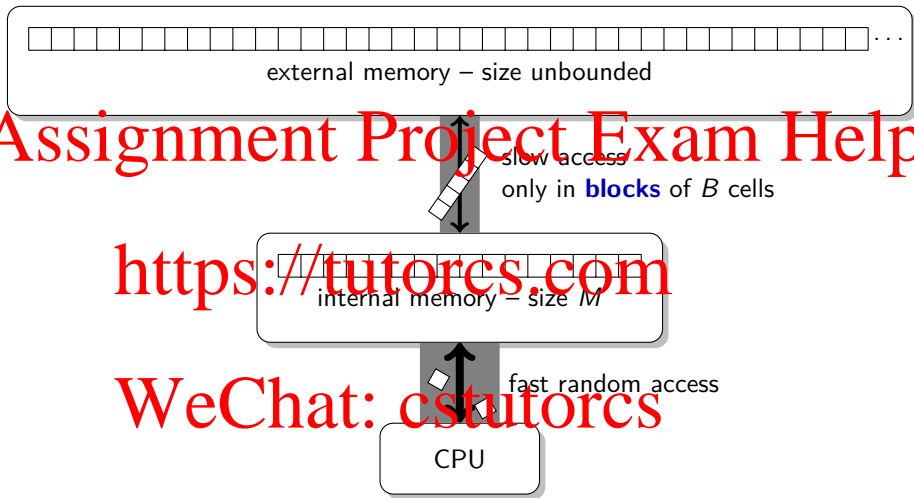
Current architectures:

- registers (very fast, very small)
- cache L1, L2 (still fast, less small)
- main memory
- disk or cloud (slow, very large)

General question: how to adapt our algorithms to take the memory hierarchy into account, avoiding transfers as much as possible?

Observation: Accessing a single location in *external memory* (e.g. hard disk) automatically loads a whole **block** (or “page”).

The External-Memory Model (EMM)



New objective: revisit all algorithms/data structures with the objective of minimizing **block transfers** (“probes”, “disk transfers”, “page loads”)

Assignment Project Exam Help

1 External Memory

- Motivation
- External sorting
- External Disks/Shares
- 2-4 Trees
- a - b -Trees
- B-Trees
- Extendable Hashing

<https://tutorcs.com>

WeChat: cstutorcs

Sorting in external memory

Recall: The sorting problem:

Given an array A of n numbers, put them into sorted order.

Now assume n is huge and A is stored in blocks in external memory.

- Heapsort was optimal in time and space in RAM model
- But: Heapsort accesses A at indices that are far apart
 - ↪ typically one block transfer per array access
 - ↪ typically $\Theta(n \log n)$ block transfers.

Can we do better?

WeChat: cstutorcs

Sorting in external memory

Recall: The sorting problem:

Given an array A of n numbers, put them into sorted order.

Now assume n is huge and A is stored in blocks in external memory.

- Heapsort was optimal in time and space in RAM model
- But: Heapsort accesses A at indices that are far apart
 - ↪ typically one block transfer per array access
 - ↪ typically $\Theta(n \log n)$ block transfers.

Can we do better?

- Mergesort adapts well to external memory. Recall algorithm:
 - ▶ Split input in half
 - ▶ Sort each half recursively \rightarrow two sorted parts
 - ▶ Merge sorted parts.

Merge

Merge(S_1, S_2)

S_1, S_2 are input streams that are in sorted order

1. $S \leftarrow$ output stream

2. while S_1 or S_2 is not empty do

```
3. if ( $S_1$  is empty)  $S.append(S_2.pop())$ 
```

```
4.     else if ( $S_2$  is empty)  $S.append(S_1.pop())$ 
```


```
5. else if ( $S_1.top() < S_2.top()$ )  $S.append(S_1.pop())$ 
```

https://tutorcs

External:

[illegible]

Internal ($B = 2, M = 8$):



Merge

$$\text{Merge}(S_1, S_2)$$

S_1, S_2 are input streams that are in sorted order

1. $S \leftarrow$ output stream

2. while S_1 or S_2 is not empty do


```
3. if ( $S_1$  is empty)  $S.append(S_2.pop())$ 
```

```
4.     else if ( $S_2$  is empty)  $S.append(S_1.pop())$ 
```


```
5. else if ( $S_1.top() < S_2.top()$ )  $S.append(S_1.pop())$ 
```

https://tutorms

External:



Internal ($B = 2, M = 8$):



Merge

Merge(S_1, S_2)

S_1, S_2 are input streams that are in sorted order

1. $S \leftarrow$ output stream

2. while S_1 or S_2 is not empty do


```
3. if ( $S_1$  is empty)  $S.append(S_2.pop())$ 
```

4. **else if** (S_2 is empty) $S.append(S_1.pop())$


```
5. else if ( $S_1.top() < S_2.top()$ )  $S.append(S_1.pop())$ 
```

```
6. else S.append($2.pop())
```

External:



Internal ($B = 2, M = 8$):




Merge

Merge(S_1, S_2)


S_1, S_2 are input streams that are in sorted order

1. $S \leftarrow$ output stream
2. while S_1 or S_2 is not empty do
3. if (S_1 is empty) $S.append(S_2.pop())$
4. else if (S_2 is empty) $S.append(S_1.pop())$
5. else if ($S_1.top() < S_2.top()$) $S.append(S_1.pop())$
6. else $S.append(S_2.pop())$

External:



Internal ($B = 2, M = 8$):




Transfer output-block to external memory when full.

Merge


```
Merge( $S_1, S_2$ )
 $S_1, S_2$  are input streams that are in sorted order
1.  $S \leftarrow$  output stream
2. while  $S_1$  or  $S_2$  is not empty do
3.   if ( $S_1$  is empty)  $S.append(S_2.pop())$ 
4.   else if ( $S_2$  is empty)  $S.append(S_1.pop())$ 
5.   else if ( $S_1.top() < S_2.top()$ )  $S.append(S_1.pop())$ 
6.   else  $S.append(S_2.pop())$ 
```

Assignment Project Exam Help
<https://tutorcs.com>

External:



Internal ($B = 2, M = 8$):



Transfer output-block to external memory when full.


Merge

Merge(S_1, S_2)


S_1, S_2 are input streams that are in sorted order

1. $S \leftarrow$ output stream
2. while S_1 or S_2 is not empty do
3. if (S_1 is empty) $S.append(S_2.pop())$
4. else if (S_2 is empty) $S.append(S_1.pop())$
5. else if ($S_1.top() < S_2.top()$) $S.append(S_1.pop())$
6. else $S.append(S_2.pop())$

External:



Internal ($B = 2, M = 8$):



Transfer output-block to external memory when full.

Load next input-block when previous is empty.


Merge

Merge(S_1, S_2)


S_1, S_2 are input streams that are in sorted order

1. $S \leftarrow$ output stream
2. while S_1 or S_2 is not empty do
3. if (S_1 is empty) $S.append(S_2.pop())$
4. else if (S_2 is empty) $S.append(S_1.pop())$
5. else if ($S_1.top() < S_2.top()$) $S.append(S_1.pop())$
6. else $S.append(S_2.pop())$

External:



Internal ($B = 2, M = 8$):



Transfer output-block to external memory when full.

Load next input-block when previous is empty.

Mergesort in external memory

- Merge takes $\Theta(n/B)$ block-transfers:

- ▶ To merge, *only* need to know first item in S_1 and S_2
 \Rightarrow only need leftmost block from S_i ($i = 1, 2$)
- ▶ Once the leftmost block of S_i is loaded, we need no other block of S_i until all its numbers have been parsed
 \Rightarrow Each block of S_i is transferred only once.
- ▶ Each block of output is transferred back only once.
- ▶ So each item is transferred twice $\Rightarrow \approx 2n/B$ block-transfers.

- Recall: Mergesort uses $O(\log_2 n)$ rounds of merging.

\Rightarrow Mergesort uses $O(n/B \cdot \log_2 n)$ block-transfers.

Better sorting in external memory

- **Observe:** We had space left in internal memory during *Merge*.


- **Idea:** We could merge d parts at once

- # Assignment Project Exam Help

- ▶ d blocks from the d input-parts
- ▶ One block for the output.

External:

<https://tutorcs.com>



5: What resources

Internal ($B = 2, M = 8$):



Better sorting in external memory


- **Observe:** We had space left in internal memory during *Merge*.

- **Idea:** We could merge d parts at once.


- Here $d \approx M/B - 1$ so that $d+1$ blocks fit into main memory.

- ▶ d blocks from the d input-parts
- ▶ One block for the output.

External:



Internal ($B = 2, M = 8$):



Better sorting in external memory

- **Observe:** We had space left in internal memory during *Merge*.


- **Idea:** We could merge d parts at once

- # Assignment Project Exam Help

- ▶ d blocks from the d input-parts
- ▶ One block for the output.


External:

<https://tutorcs.com>



5:  WeChat: cstutorcs

Internal ($B = 2, M = 8$):



Better sorting in external memory


- **Observe:** We had space left in internal memory during *Merge*.

- **Idea:** We could merge d parts at once.


- Here $d \approx M/B - 1$ so that $d+1$ blocks fit into main memory.

- ▶ d blocks from the d input-parts
- ▶ One block for the output.

External:



Internal ($B = 2, M = 8$):



Better sorting in external memory


- **Observe:** We had space left in internal memory during *Merge*.

- **Idea:** We could merge d parts at once.


- Here $d \approx M/B - 1$ so that $d+1$ blocks fit into main memory.

- ▶ d blocks from the d input-parts
- ▶ One block for the output.

External:



Internal ($B = 2, M = 8$):



← S_1 → ← S_2 → ← S_3 → ← output →

Better sorting in external memory


- **Observe:** We had space left in internal memory during *Merge*.

- **Idea:** We could merge d parts at once.


- Here $d \approx M/B - 1$ so that $d+1$ blocks fit into main memory.

- ▶ d blocks from the d input-parts
- ▶ One block for the output.

External:



Internal ($B = 2, M = 8$):



Better sorting in external memory


- **Observe:** We had space left in internal memory during *Merge*.

- **Idea:** We could merge d parts at once.


- Here $d \approx M/B - 1$ so that $d+1$ blocks fit into main memory.

- ▶ d blocks from the d input-parts
- ▶ One block for the output.

External:



Internal ($B = 2, M = 8$):



← S_1 → ← S_2 → ← S_3 → ← output →

d-way merge

- When consider the *internal run-time*, d -way merge is fastest when using a *min-oriented* priority queue.

Assignment Project Exam Help

d-way-merge (S_1, \dots, S_d)

S_1, \dots, S_d are input streams that are in sorted order

1. $P \leftarrow$ empty min-priority queue
2. $S \leftarrow$ output stream
3. **for** $i \leftarrow 1$ **to** d **do** $P.insert((S_i.top(), i))$
 // each item in P keeps track of its input-stream
4. **while** P is not empty **do**
5. $(x, i) \leftarrow P.deleteMin()$
6. $S.append(S_i.pop())$
7. **if** S_i is not empty **do**
8. $P.insert((S_i.top(), i))$

- The run-time for *d-way-merge* is $O(n \log d)$.
- The number of *block transfers* is again $O(n/B)$.

d-way mergesort

Idea: Apply Mergesort but split into many parts.

- Set $d \approx M/B - 1$
- Split the input into d approximately equal parts S_1, \dots, S_d .
- Recursively sort S_1, \dots, S_d .
- *d-way-merge*(S_1, \dots, S_d)

Run-time: $T(n) = dT(n/d) + O(n \log d)$

- The recursion-depth is $O(\log_d n)$
- On each level, we do $O(n \log d)$ work.
- Total time: $O(\log_d n \cdot n \log d) = O(n \log n)$
- In the RAM model (internal memory), this is no better than regular Mergesort.

Block-transfers: We have $O(n/B)$ block-transfers per level.

- Total # block-transfers: $O(\log_d n \cdot (n/B))$.

Mergesort with external memory

Total # block transfers with d -way mergesort: $O(\log_d(n) \cdot (n/B))$

Assignment Project Exam Help

One can prove lower bounds in the external memory model:

Any comparison-based sorting algorithm requires

$\Omega\left(\frac{n}{B} \log_{M/B}\left(\frac{n}{B}\right)\right)$ *block transfers*

<https://tutorcs.com>

(The proof is beyond the scope of the course.)

WeChat: cstutorcs

Mergesort with external memory

Total # block transfers with d -way mergesort: $O(\log_d(n) \cdot (n/B))$

Assignment Project Exam Help

One can prove lower bounds in the external memory model:

Any comparison-based sorting algorithm requires

$\Omega(\frac{n}{B} \log_{M/B}(\frac{n}{B}))$ block transfers

<https://tutorcs.com>

(The proof is beyond the scope of the course.)

WeChat: estutorcs

- Recall: We can use $d \in \Theta(M/B)$ for Mergesort.
- d -way mergesort is optimal (up to constant factors)!

Assignment Project Exam Help

1 External Memory

- Motivation
- External sorting
- External Dictionaries
- 2-4 Trees
- a - b -Trees
- B-Trees
- Extendible Hashing

<https://tutorcs.com>

WeChat: cstutorcs

Dictionaries in external memory

Recall: Dictionaries store n KVPs and support *search*, *insert* and *delete*.

Assignment Project Exam Help

- **Recall:** AVL-trees were optimal in time and space in RAM model
- But: Inserts happen at varying locations of the tree.
 - ↪ nearby nodes are unlikely to be on the same block
 - ↪ typically $\Theta(\log n)$ block transfers per operation


Better solution: design a tree-structure that *guarantees* that many nodes on search-paths are within one block.

https://tutorcs.com

WeChat: cstutorcs

Idealized structure

block of external memory




Idea: Store subtrees in one block of memory.

- If block can hold subtree of size $b-1$, then block covers height $\log b$

\Rightarrow Search-path hits $\frac{\Theta(\log n)}{\log b}$ blocks $\Rightarrow \Theta(\log_b n)$ block-transfers

Idealized structure

block of external memory



Idea: Store subtrees in one block of memory.

- If block can hold subtree of size $b-1$, then block covers height $\log b$
- ⇒ Search-path hits $\frac{\Theta(\log n)}{\log b}$ blocks ⇒ $\Theta(\log_b n)$ block-transfers
- View block as one node of a *multiway-tree* ($b-1$ KVPs, b children)
- To allow *insert/delete*, we permit varying numbers of KVPs in nodes

Assignment Project Exam Help

1 External Memory

- Motivation
- External sorting
- External Disks/Shares
- 2-4 Trees
- a - b -Trees
- B-Trees
- Extendable Hashing

<https://tutorcs.com>

WeChat: cstutorcs

2-4 Trees


We will first explore this idea for $b = 4$. The resulting search tree is also useful for internal memory considerations.

Structural properties:


- Every node is either
 - ▶ 1-node: *one KVP* and *two subtrees* (possibly empty), or
 - ▶ 2-node: *two KVPs* and *three subtrees* (possibly empty), or
 - ▶ 3-node: *three KVPs* and *four subtrees* (possibly empty).
- All empty subtrees are at the same level.

Summary: Height-balance is strictly enforced, but allow 3 types of nodes!

Order property: The keys at a node are between the keys in the subtrees.



Assignment Project Exam Help



2-4 Tree operations

search: The order-property determines the subtree to search in.

24Tree::search($k, v \leftarrow \text{root}, p \leftarrow \text{NIL}$)


k : key to search, v : node where we search, p : parent of v

1. **if** v represents empty subtree
2. **return** "not found, would be in p "
3. Let $\langle T_0, k_1, \dots, k_d, T_d \rangle$ be key-subtree list at v
4. **if** $k \geq k_1$
5. $i \leftarrow$ maximal index such that $k_i \leq k$
6. **if** $k_i = k$
7. **return** "at i th key in v "
8. **else** *24Tree::search*(k, T_i, v)
9. **else** *24Tree::search*(k, T_0, v)

2-4 Tree operations

24Tree::insert(k)


1. $v \leftarrow \text{24Tree::search}(k)$ // leaf where k should be
2. Add k and an empty subtree in key-subtree-list of v
3. **while** v has 4 keys (**overflow** \rightsquigarrow **node split**)
4. Let $\langle T_0, k_1, \dots, k_4, T_4 \rangle$ be key-subtree list at v
5. **if** (v has no parent) create a parent of v without KVPs
6. $p \leftarrow$ parent of v
7. $v' \leftarrow$ new node with keys k_1, k_2 and subtrees T_0, T_1, T_2
8. $v'' \leftarrow$ new node with key k_4 and subtrees T_3, T_4
9. Replace $\langle v \rangle$ by $\langle v', k_3, v'' \rangle$ in key-subtree-list of p
10. $v \leftarrow p$



Example: Insertion in a 2-4 tree

Example: *insert*(17)


Assignment Project Exam Help



Example: Insertion in a 2-4 tree

Example: *insert*(17)


Assignment Project Exam Help



Example: Insertion in a 2-4 tree

Example: *insert*(17)

Assignment Project Exam Help




Example: Insertion in a 2-4 tree

Example: *insert*(17)

Assignment Project Exam Help


<https://tutorcs.com>

WeChat: cstutorcs



Example: Insertion in a 2-4 tree

Example: *insert*(17)



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs


Deletion from a 2-4 Tree

24Tree::delete(k)

1. $w \leftarrow \text{24Tree::search}(k)$ // node containing k
2. $v \leftarrow$ leaf containing predecessor or successor k' of k
Replace k by k' in w , delete k' and empty subtree in v
3. **while** v has 0 keys (**underflow**)
4. **if** v is the root, delete it and **break**
5. $p \leftarrow$ parent of v
6. **if** v has a sibling u with 2 or more keys (**transfer/rotate**)
7. **if** u is right sibling
8. Replace key k in p by $u.k_1$
9. Remove $\langle u.T_0, u.k_1 \rangle$ from u , append $\langle k, u.T_0 \rangle$ to v
10. **else** ... // symmetrically with left sibling
11. **else** (**merge & repeat**)
12. **if** v has right sibling u
13. $v' \leftarrow$ new node with list $\langle v.T_0, k, u.T_0, u.k_1, u.T_1 \rangle$
14. replace $\langle v, k, u \rangle$ by $\langle v' \rangle$ in p
15. $v \leftarrow p$
16. **else** ... // symmetrically with left sibling
- 17.


2-4 Tree Deletion

Example: *delete*(43)




2-4 Tree Deletion

Example: *delete*(43)



2-4 Tree Deletion

Example: *delete*(43)



2-4 Tree Deletion

Example: *delete*(19)




2-4 Tree Deletion

Example: *delete*(19)

Assignment Project Exam Help


<https://tutorcs.com>

WeChat: cstutorcs




2-4 Tree Deletion

Example: *delete*(19)




2-4 Tree Deletion

Example: *delete*(42)




2-4 Tree Deletion

Example: *delete*(42)



2-4 Tree Deletion

Example: *delete*(42)




2-4 Tree Deletion

Example: *delete*(42)

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs




2-4 Tree Deletion

Example: *delete*(42)

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Assignment Project Exam Help

1 External Memory

- Motivation
- External sorting
- External Disks/Shares
- 2-4 Trees
- ***a-b*-Trees**
- B-Trees
- Extendable Hashing

<https://tutorcs.com>

WeChat: cstutorcs

a - b -Trees

A 2-4 tree is an a - b -tree for $a = 2$ and $b = 4$.

An a - b -tree satisfies:


- Each node has at least a subtrees, unless it is the root.
The root has at least 2 subtrees.
- Each node has at most b subtrees.
- If a node has k subtrees, then it stores $k - 1$ key-value pairs (KVPs).
- Empty subtrees are at the same level.
- The keys in the node are between the keys in the corresponding subtrees.

Requirement: $a \leq \lceil b/2 \rceil$.

search, *insert*, *delete* then work just like for 2-4 trees, after re-defining underflow/overflow to consider the above constraints.

a-b-tree example

A 3-6 tree




Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs


a-b-tree insertion

insert(55):



a-b-tree insertion

insert(55):



Height of an a - b -tree

Recall: n = numbers of KVPs (*not* the number of nodes)

What is smallest possible number of KVPs in an a - b -tree of height- h ?

Assignment Project Exam Help

| Level | Nodes \geq | Links/node \geq | KVP/node \geq | KVPs on level \geq |
|-------|--------------|-------------------|-----------------|----------------------|
| 0 | 1 | 2 | 1 | 1 |
| 1 | 2 | a | $a - 1$ | $2(a - 1)$ |
| 2 | $2a$ | a | $a - 1$ | $2a(a - 1)$ |
| 3 | $2a^2$ | a | $a - 1$ | $2a^2(a - 1)$ |
| ... | ... | ... | ... | ... |
| h | $2a^{h-1}$ | a | $a - 1$ | $2a^{h-1}(a - 1)$ |

<https://tutorcs.com>

WeChat: cstutorcs

$$\text{Total: } n = \# \text{KVPs} \geq 1 + 2(a - 1) \sum_{i=0}^{h-1} a^i = 2a^h - 1$$

Therefore the height of an a - b -tree is $O(\log_a(n)) = O(\log n / \log a)$.

a-b-trees as implementations of dictionaries

Analysis (if entire a - b -tree is stored in internal memory):

- *search*, *insert*, and *delete* each requires visiting $\Theta(\text{height})$ nodes

- Height is $O(\log n / \log a)$

- Recall: $a \leq \lceil b/2 \rceil$ required for *insert* and *delete*

\Rightarrow choose $a = \lceil b/2 \rceil$ to minimize the height.

- Work at a node can be done in $O(\log b)$ time.

$$\text{Total cost: } O\left(\frac{\log n}{\log a} \cdot (\log b)\right) = O\left(\log n \cdot \frac{\log b}{\log b - 1}\right) = O(\log n)$$

This is no better than AVL-trees.

(Though 2-4-trees are faster than AVL-trees in practice, especially when converted to binary search trees called *red-black trees*. No details.)

The main motivation for a - b -trees is *external memory*.

Assignment Project Exam Help

1 External Memory

- Motivation
- External sorting
- External Disks/Shares
- 2-4 Trees
- a - b -Trees
- B-Trees
- Extensible Hashing

<https://tutorcs.com>

WeChat: cstutorcs

B-trees


A **B-tree** is an *a-b*-tree tailored to the external memory model.

- Every node is one block of memory (of size B).


- b is chosen maximally such that a node with $b-1$ KVPs (hence $b-1$ value-references and b subtree-references) fits into a block.

b is called the **order** of the B -tree. Typically $b \in \Theta(B)$.

- a is set to be $\lceil b/2 \rceil$ as before.



B-tree analysis



- *search*, *insert*, and *delete* each requires visiting $\Theta(\text{height})$ nodes
- Work within a node is done in internal memory \Rightarrow no block-transfer.
- The height is $\Theta(\log_a n) = \Theta(\log_B n)$ (presuming $a = \lceil b/2 \rceil \in \Theta(B)$)

So all operations require $\Theta(\log_B n)$ **block transfers**.

This is asymptotically optimal.

B-tree variations

For practical purposes, some variations are better:

- B-trees with **pre-emptive splitting/merging**:

- ▶ During search for insert, split *any* node close to overflow.
- ▶ During search for delete, merge *any* node close to underflow.

↪ can insert/delete at leaf and stop, this halves block transfers.

- **B⁺-trees**: All KVPs are in leaves.

- ▶ Interior nodes store keys to guide search-path, but no values.
- ▶ Leaves omit references to empty subtrees.

↪ bigger order, hence smaller height

- **Cache-oblivious** trees: What if we do not know B ?

- ▶ Build an a - b -tree with $b \approx \sqrt{n}$
- ▶ Each node stores its KVPs again as cache-oblivious tree.

↪ achieves $O(\log_B(n))$ block transfers *without* knowing B .

Assignment Project Exam Help

1 External Memory

- Motivation
- External sorting
- External Disks/Streams
- 2-4 Trees
- a - b -Trees
- B-Trees
- Extendible Hashing

<https://tutorcs.com>

WeChat: cstutorcs

Dictionaries for Integers in External Memory

- Recall: Direct Addressing allowed for $O(1)$ insert and delete if keys are integers in $\{0, \dots, U-1\}$

- If keys are too big, hashing was used to map keys to (smaller) integers.


- This does not adapt well to external memory.

- ▶ Most hash strategies access many blocks (probe sequence is scattered)
- ▶ Even those that do not (Linear Probing, Cuckoo hashing) need to re-hash to keep α small.
- ▶ And re-hashing must load *all* blocks

- New Idea: Store tree of links to blocks of integers.

(This is also called **extendible hashing**, because its primary use is for dictionaries that store integers that result from hashing.)

Tries of blocks – Overview



Assumption: We store integers in $\{0, 1, \dots, 2^L - 1\}$.

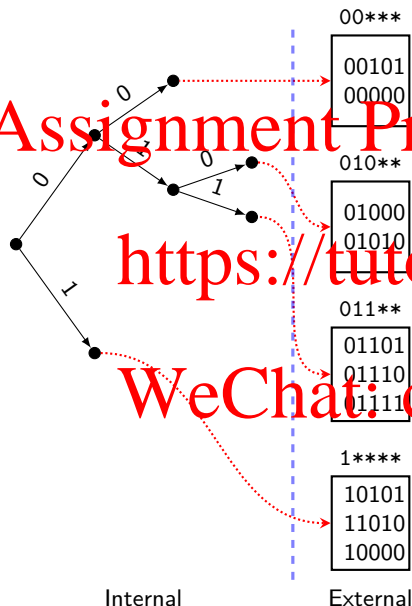
Interpret all integers as bitstrings of length L .

Build trie D (the **directory**) of integers in internal memory.

Stop splitting in trie when remaining items fit in one block.

Each leaf of D refers to block of external memory that stores the items.

External hashing with tries – Details



search(k): Search for $(k)_2$ in D until we reach leaf ℓ . Load block at ℓ and search in it.
1 block transfer.

insert(k): Search for k and load block, then insert k . If this exceeds block-capacity, split at trie-node and split blocks (possibly repeatedly).
Typically 2 block transfers.

delete(k): Search for k and load block, then use lazy deletion.
Optional: combine underfull blocks.
2 block transfers.

Extendible hashing: Insert

ExtendibleHashing::insert(k)

k : integer in $\{0, \dots, 2^L - 1\}$

1. Convert k to length- L bitstring w
2. $\ell \leftarrow \text{Trie}::\text{search}(D, w)$ // leaf where w would be
3. $d \leftarrow \text{depth of } \ell \text{ in } D$
4. transfer block P that ℓ refers to
5. **while** P has no room for additional items
6. Split P into two blocks P_0 and P_1 by $(d+1)^{\text{st}}$ digit
7. Create two children ℓ_0 and ℓ_1 of ℓ , linked to P_0 and P_1
8. $d \leftarrow d+1$, $\ell \leftarrow \ell_{w[d]}$, $P \leftarrow P_{w[d]}$
9. insert w into P


Note: This may create empty blocks, but this should be rare.

insert(10110)

Assignment Project Exam Help

<https://tutorcs.com>


WeChat: cstutorcs



Extendible hashing: saving space

We can save links (hence space in internal memory) with two tricks:

- Expand the trie so that all leaves have the same **global depth** d .
- Store **only** the leaves, and in an array D of size 2^d .
- Operations work as before if each block stores its **local depth**, i.e. the depth of the original trie-node that referred to it.



Extendible hashing discussion

- Hashing collisions (= duplicate keys) may happen, but are resolved within the block and do not affect the block transfers.

If more items collide than can fit into a block we use other strategies (chaining, open addressing). This should be exceedingly rare.

- Directory is much smaller than total number of stored keys

→ should fit in internal memory.

If it does not, then strategies similar to B-trees can be applied.

- Only 1 or 2 block transfers expected for *any* operation.

- To make more space, we only add one block.

Rarely change the size of the directory.

Never have to move all items. (in contrast to re-hashing!)

- Space usage is not too inefficient: one can show that under uniform distribution assumption each block is expected to be 69% full.